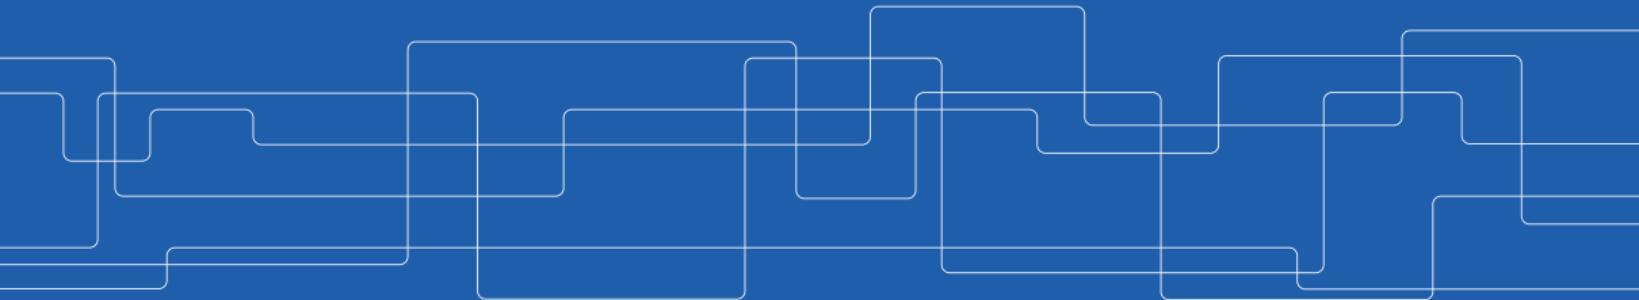




Distributed Deep Learning

Amir H. Payberah
payberah@kth.se
2021-12-08



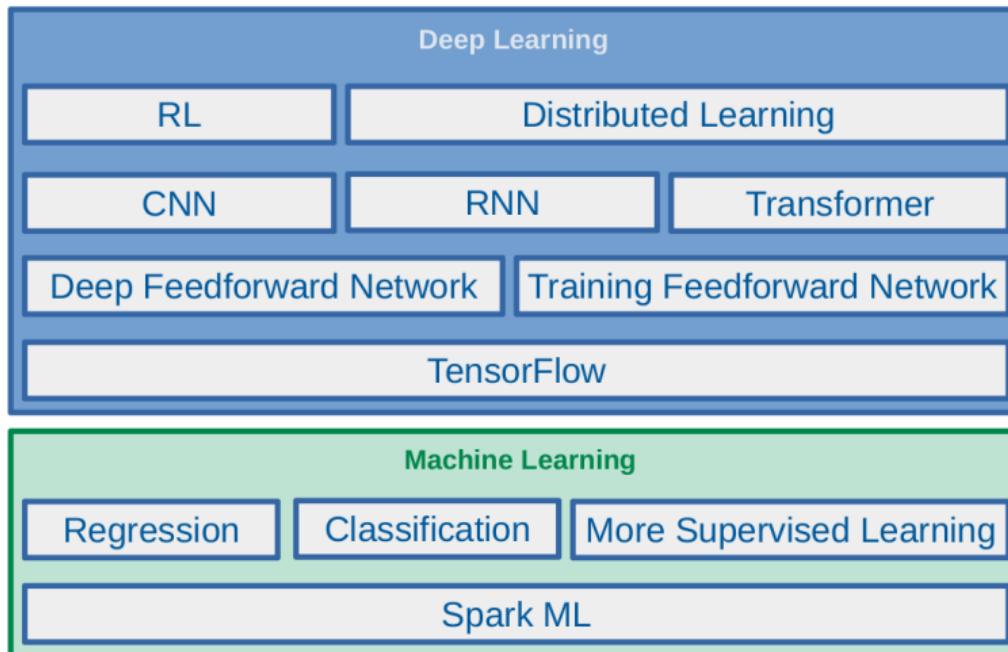


The Course Web Page

<https://id2223kth.github.io>
<https://tinyurl.com/6s5jy46a>

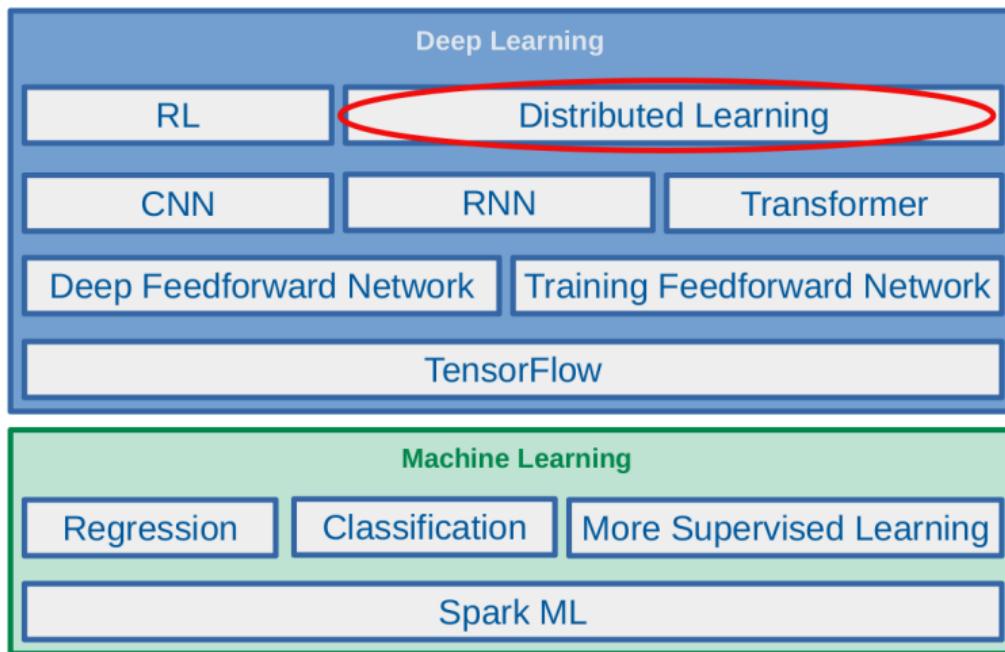


Where Are We?





Where Are We?

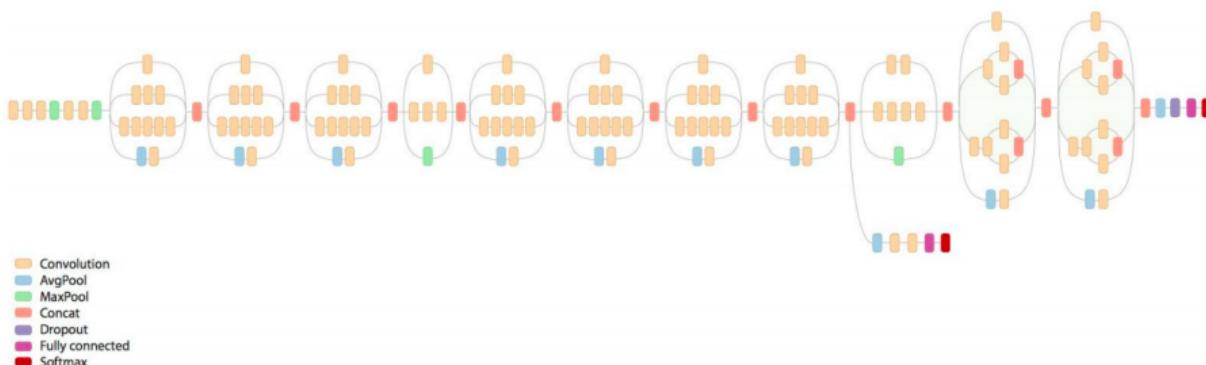




What is the problem?

Training Deep Neural Networks

- ▶ Computationally intensive
- ▶ Time consuming



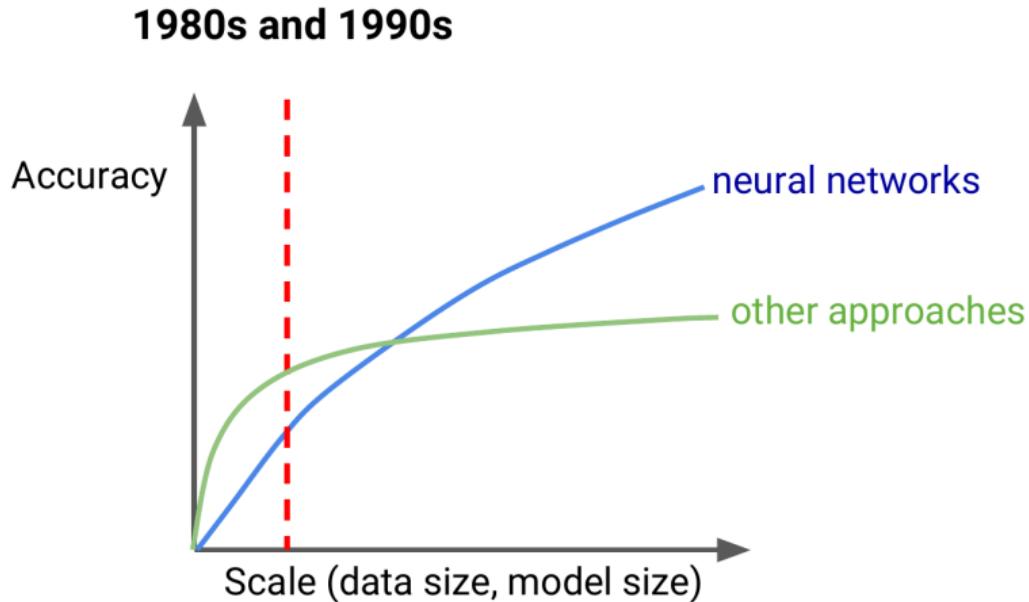
[<https://cloud.google.com/tpu/docs/images/inceptionv3onc--oview.png>]

Why?

- ▶ Massive amount of training dataset
- ▶ Large number of parameters

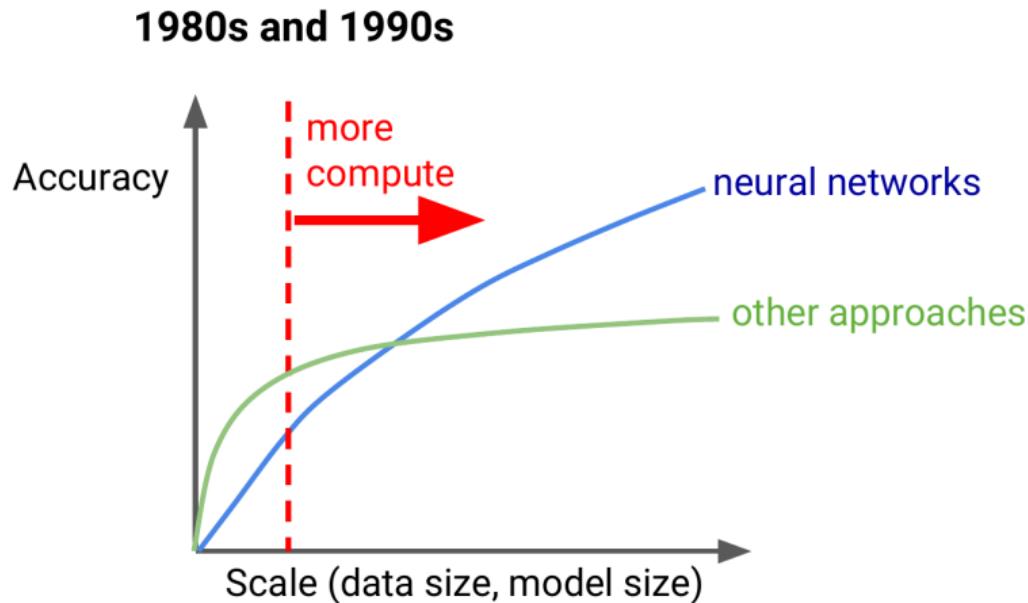


Accuracy vs. Data/Model Size



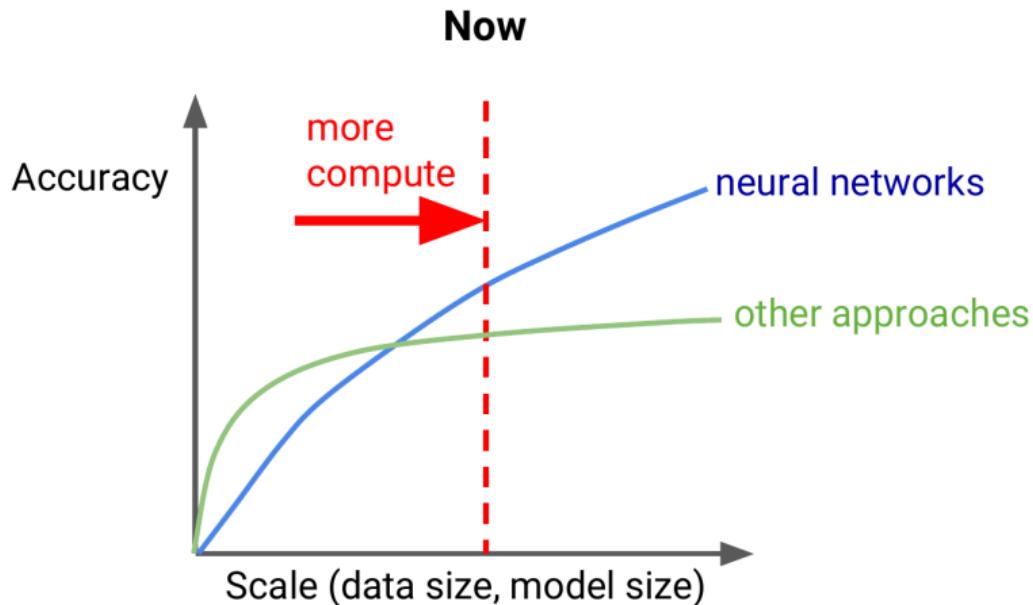
[Jeff Dean at AI Frontiers: Trends and Developments in Deep Learning Research]

Accuracy vs. Data/Model Size



[Jeff Dean at AI Frontiers: Trends and Developments in Deep Learning Research]

Accuracy vs. Data/Model Size



[Jeff Dean at AI Frontiers: Trends and Developments in Deep Learning Research]

Scalability



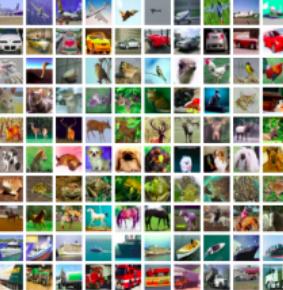


Fundamentals of Machine Learning



Training Dataset

- ▶ E.g., tabular data, image, text, etc.



Entities

Society and Culture Science and Mathematics Health Education and Reference Computers and Internet Sports
Business and Finance Entertainment and Music Family and Relationships Politics and Government

does anyone here play habbohotel and want 2 be friends? Answer: No on the first part and maybe on the second part. I got to think it over first.

Family and Relationships

Date	Cost	Actions	Offsite conversions	Impressions	Clicks
2017-04-04	29.44	461	4	5655	477
2017-04-03	74.08	1331	16	18170	1340
2017-04-02	76.09	1349	12	16877	1357
2017-04-01	76.79	1382	8	19757	1378
2017-03-31	77.28	1141	21	18598	1116
2017-03-30	68.62	1065	18	14847	1046
2017-03-29	64.9	1111	25	13994	1094
2017-03-28	65.12	1137	12	15952	1145
2017-03-27	66.98	1185	7	17970	1190
2017-03-26	64.94	1118	5	14410	1116
2017-03-25	66.3	1208	6	15123	1204
2017-03-24	67.38	1143		15296	1159
2017-03-23	65.59	1147	13	14972	1143
2017-03-22	68.19	1129	4	17959	1116
2017-03-21	64.78	1081		25810	1059

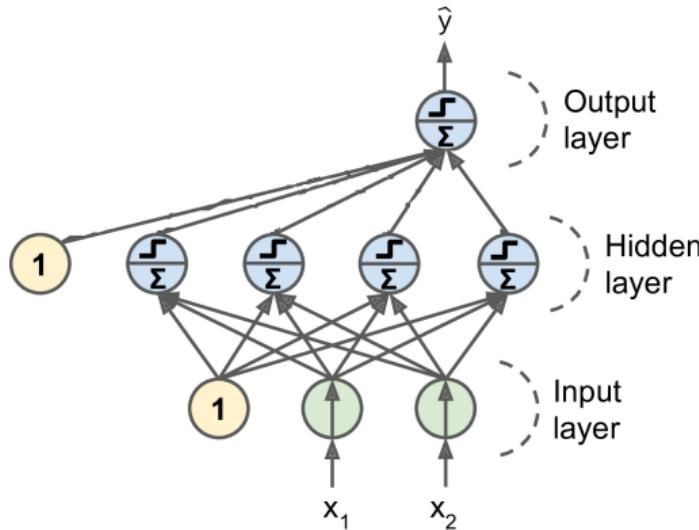


Model

- ▶ E.g., linear models, neural networks, etc.

Model

- ▶ E.g., linear models, neural networks, etc.
- ▶ $\hat{y} = f_w(x)$





Loss function

- ▶ How good \hat{y} is able to predict the expected outcome y .

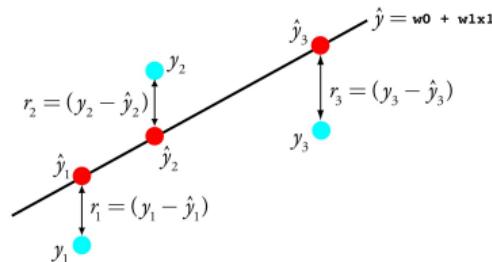


Loss function

- ▶ How good \hat{y} is able to predict the expected outcome y .
- ▶ $J(w) = \sum_{i=1}^m l(y_i, \hat{y}_i)$

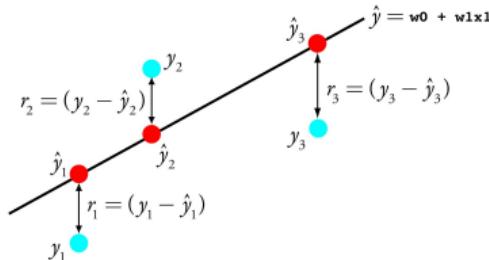
Loss function

- ▶ How good \hat{y} is able to predict the expected outcome y .
- ▶ $J(w) = \sum_{i=1}^m l(y_i, \hat{y}_i)$



Loss function

- ▶ How good \hat{y} is able to predict the expected outcome y .
- ▶ $J(w) = \sum_{i=1}^m l(y_i, \hat{y}_i)$



- ▶ E.g., $J(w) = \frac{1}{m} \sum_{i=1}^m (y_i - \hat{y}_i)^2$



Objective

- ▶ Minimize the loss function

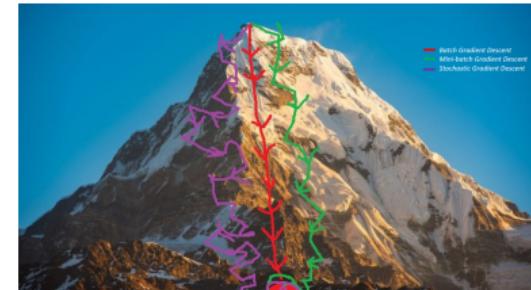


Objective

- ▶ Minimize the loss function
- ▶ $\arg \min_w J(w)$
- ▶ $J(w) = \sum_{i=1}^m l(y_i, \hat{y}_i)$

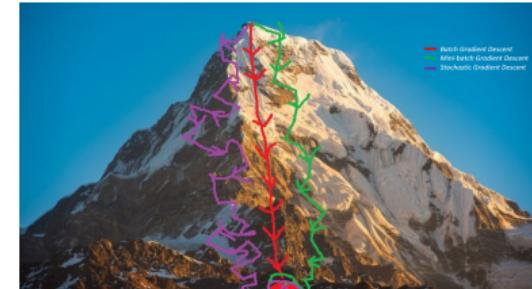
Training

► $J(w) = \sum_{i=1}^m l(y_i, \hat{y}_i)$



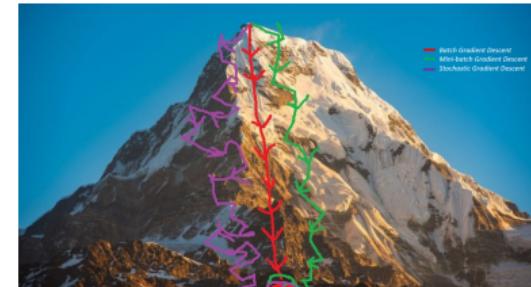
Training

- ▶ $J(w) = \sum_{i=1}^m l(y_i, \hat{y}_i)$
- ▶ Gradient descent, i.e., $w := w - \eta \nabla J(w)$



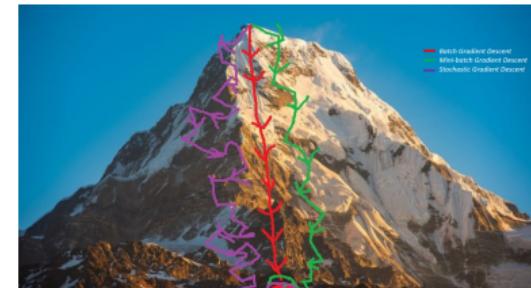
Training

- ▶ $J(w) = \sum_{i=1}^m l(y_i, \hat{y}_i)$
- ▶ Gradient descent, i.e., $w := w - \eta \nabla J(w)$
- ▶ Stochastic gradient descent, i.e., $w := w - \eta \tilde{g} J(w)$
 - \tilde{g} : gradient at a randomly chosen point.



Training

- ▶ $J(w) = \sum_{i=1}^m l(y_i, \hat{y}_i)$
- ▶ Gradient descent, i.e., $w := w - \eta \nabla J(w)$
- ▶ Stochastic gradient descent, i.e., $w := w - \eta \tilde{g} J(w)$
 - \tilde{g} : gradient at a **randomly** chosen point.
- ▶ Mini-batch gradient descent, i.e., $w := w - \eta \tilde{g}_B J(w)$
 - \tilde{g} : gradient with respect to a set of B **randomly** chosen points.





Let's Scale the Learning



Scalable Training

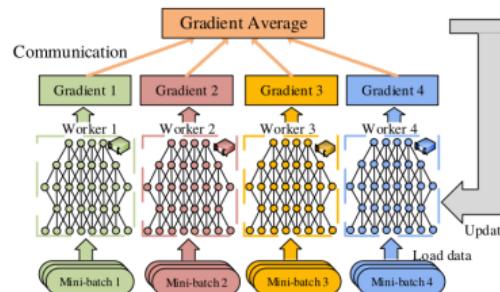
- ▶ Data parallelism
- ▶ Model parallelism



Data Parallelism

Data Parallelization (1/4)

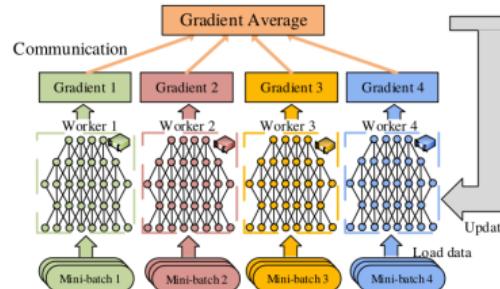
- ▶ Replicate a **whole model** on **every device**.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Data Parallelization (1/4)

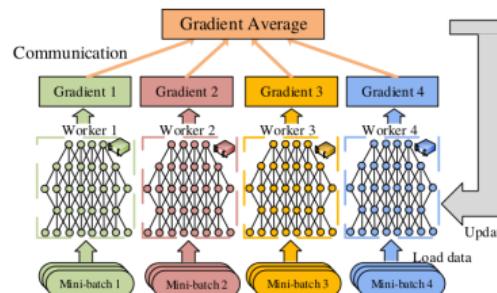
- ▶ Replicate a **whole model** on **every device**.
- ▶ Train **all replicas simultaneously**, using a **different mini-batch** for each.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Data Parallelization (2/4)

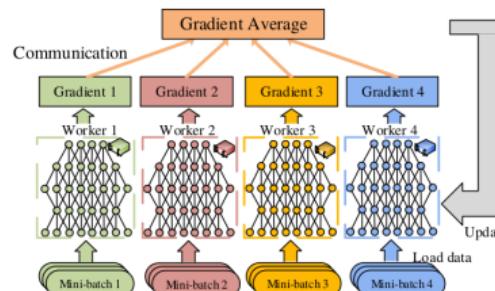
- ▶ k devices



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Data Parallelization (2/4)

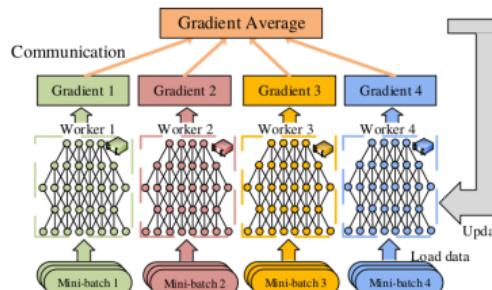
- ▶ k devices
- ▶ $J_j(w) = \sum_{i=1}^{b_j} l(y_i, \hat{y}_i), \forall j = 1, 2, \dots, k$



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Data Parallelization (2/4)

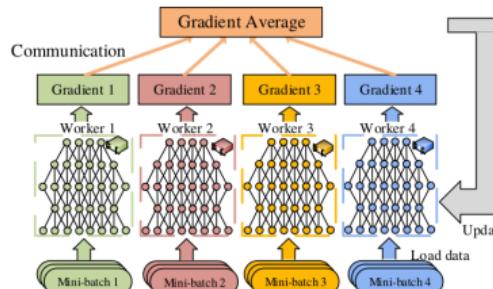
- ▶ k devices
- ▶ $J_j(w) = \sum_{i=1}^{b_j} l(y_i, \hat{y}_i), \forall j = 1, 2, \dots, k$
- ▶ $\tilde{g}_B J_j(w)$: gradient of $J_j(w)$ with respect to a set of B randomly chosen points at device j .



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Data Parallelization (2/4)

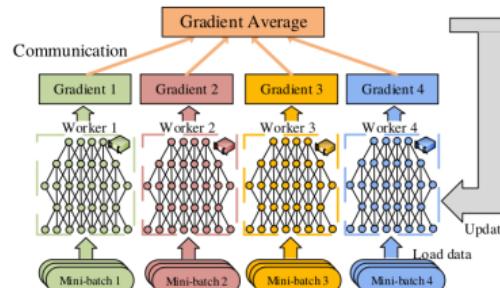
- ▶ k devices
- ▶ $J_j(w) = \sum_{i=1}^{b_j} l(y_i, \hat{y}_i), \forall j = 1, 2, \dots, k$
- ▶ $\tilde{g}_B J_j(w)$: gradient of $J_j(w)$ with respect to a set of B randomly chosen points at device j .
- ▶ Compute $\tilde{g}_B J_j(w)$ on each device j .



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Data Parallelization (3/4)

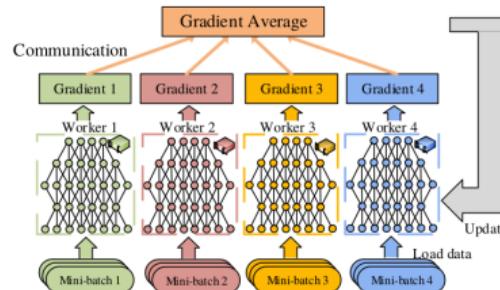
- ▶ Compute the **mean of the gradients**.
- ▶ $\tilde{g}_B J(w) = \frac{1}{k} \sum_{j=1}^k \tilde{g}_B J_j(w)$



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Data Parallelization (4/4)

- ▶ Update the model.
- ▶ $w := w - \eta \tilde{g}_B J(w)$



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]



Data Parallelization Design Issues

- ▶ The **aggregation** algorithm
- ▶ Communication **synchronization** and frequency
- ▶ Communication **compression**



The Aggregation Algorithm



The Aggregation Algorithm

- ▶ How to aggregate gradients (compute the mean of the gradients)?



The Aggregation Algorithm

- ▶ How to aggregate gradients (compute the mean of the gradients)?
- ▶ Centralized - parameter server



The Aggregation Algorithm

- ▶ How to aggregate gradients (compute the mean of the gradients)?
- ▶ Centralized - parameter server
- ▶ Decentralized - all-reduce



The Aggregation Algorithm

- ▶ How to aggregate gradients (compute the mean of the gradients)?
- ▶ Centralized - parameter server
- ▶ Decentralized - all-reduce
- ▶ Decentralized - gossip

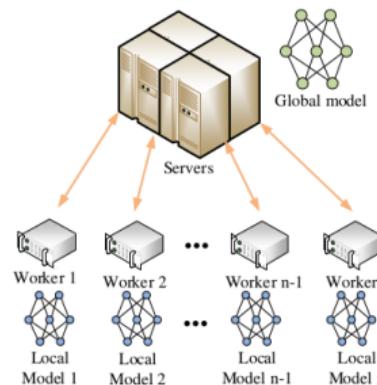


Aggregation - Centralized - Parameter Server

- ▶ Store the model parameters **outside of the workers**.

Aggregation - Centralized - Parameter Server

- ▶ Store the model parameters **outside of the workers**.
- ▶ **Workers** periodically report their **computed parameters** or **parameter updates** to a (set of) **parameter server(s) (PSs)**.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

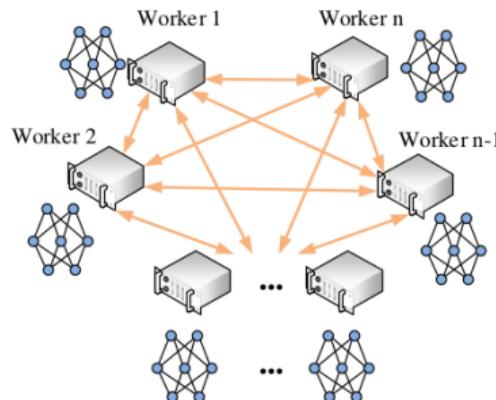


Aggregation - Distributed - All-Reduce

- ▶ Mirror all the model parameters across all workers (no PS).

Aggregation - Distributed - All-Reduce

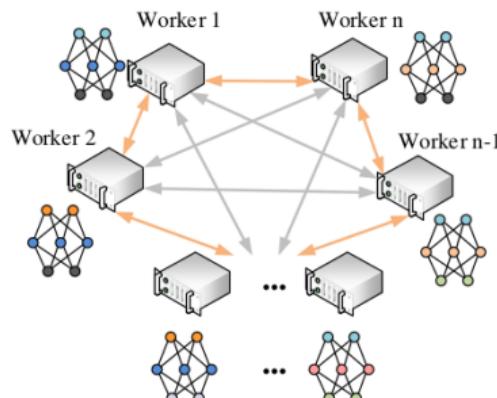
- ▶ Mirror all the model **parameters** across all workers (no PS).
- ▶ Workers **exchange** parameter updates **directly** via an **allreduce** operation.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Aggregation - Distributed - Gossip

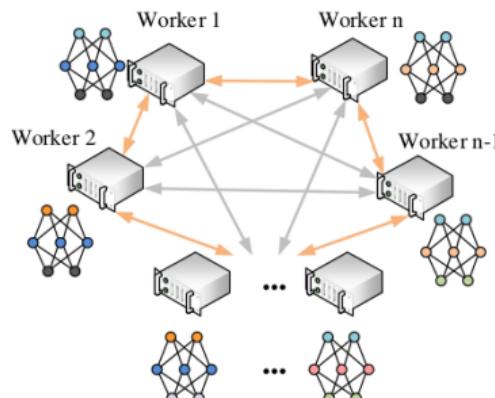
- ▶ No PS, and no global model.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Aggregation - Distributed - Gossip

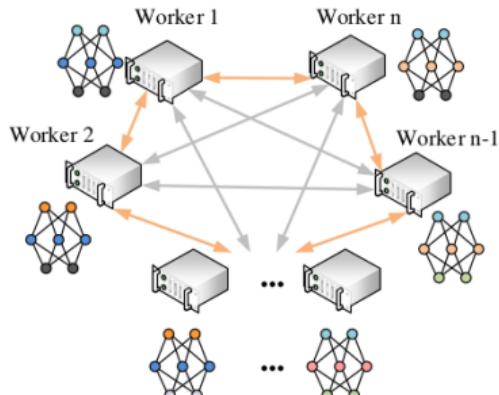
- ▶ No PS, and no global model.
- ▶ Every worker communicates updates with their neighbors.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Aggregation - Distributed - Gossip

- ▶ No PS, and no global model.
- ▶ Every worker communicates updates with their neighbors.
- ▶ The consistency of parameters across all workers only at the end of the algorithm.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]



Reduce and AllReduce (1/2)

- ▶ **Reduce**: reducing a **set of numbers** into a **smaller set of numbers** via a function.

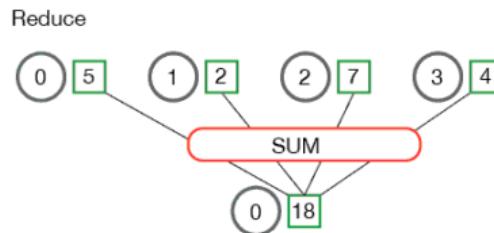


Reduce and AllReduce (1/2)

- ▶ **Reduce**: reducing a **set of numbers** into a **smaller set of numbers** via a function.
- ▶ E.g., `sum([1, 2, 3, 4, 5]) = 15`

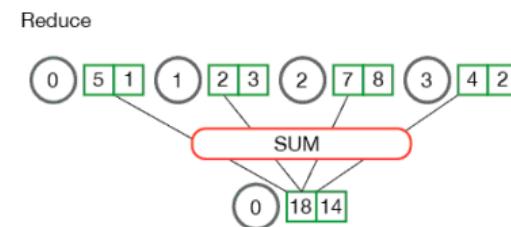
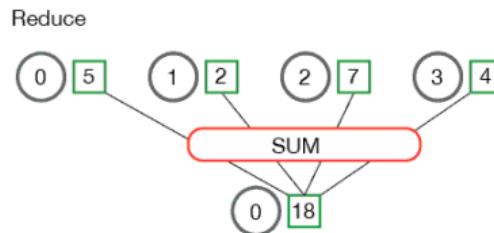
Reduce and AllReduce (1/2)

- ▶ **Reduce**: reducing a **set of numbers** into a **smaller set of numbers** via a function.
- ▶ E.g., `sum([1, 2, 3, 4, 5]) = 15`
- ▶ Reduce takes an **array of input** elements on each process and returns an **array of output** elements to the **root process**.



Reduce and AllReduce (1/2)

- ▶ **Reduce**: reducing a **set of numbers** into a **smaller set of numbers** via a function.
- ▶ E.g., `sum([1, 2, 3, 4, 5]) = 15`
- ▶ Reduce takes an **array of input** elements on each process and returns an **array of output** elements to the **root process**.



[<https://mpitutorial.com/tutorials/mpi-reduce-and-allreduce>]

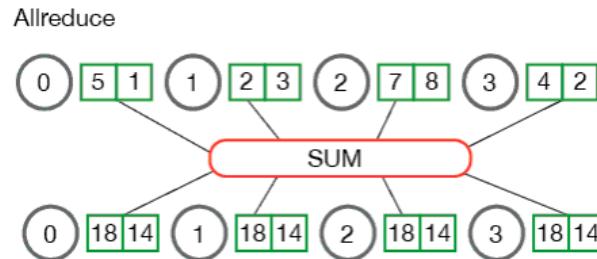


Reduce and AllReduce (2/2)

- ▶ AllReduce stores reduced results across all processes rather than the root process.

Reduce and AllReduce (2/2)

- ▶ AllReduce stores reduced results across all processes rather than the root process.



[<https://mpitutorial.com/tutorials/mpi-reduce-and-allreduce>]

AllReduce Example

Initial state

Worker A
17 11 1 9

Worker B
5 13 23 14

Worker C
3 6 10 8

Worker D
12 7 2 12



After AllReduce operation

Worker A
37 37 36 43

Worker B
37 37 36 43

Worker C
37 37 36 43

Worker D
37 37 36 43

[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

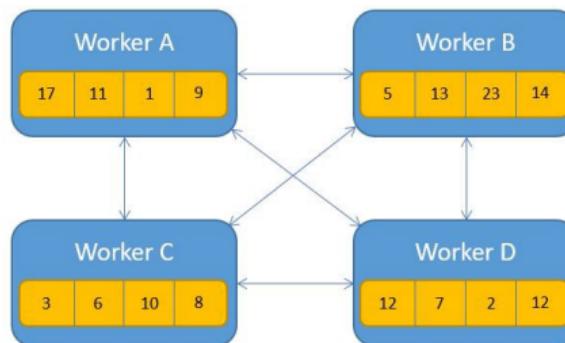


AllReduce Implementation

- ▶ All-to-all allreduce
- ▶ Master-worker allreduce
- ▶ Tree allreduce
- ▶ Round-robin allreduce
- ▶ Butterfly allreduce
- ▶ Ring allreduce

AllReduce Implementation - All-to-All AllReduce

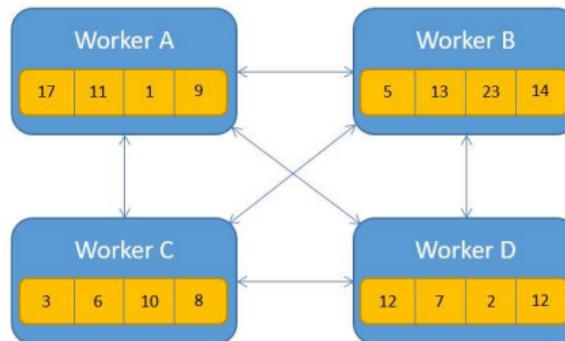
- ▶ Send the array of data to each other.
- ▶ Apply the reduction operation on each process.



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - All-to-All AllReduce

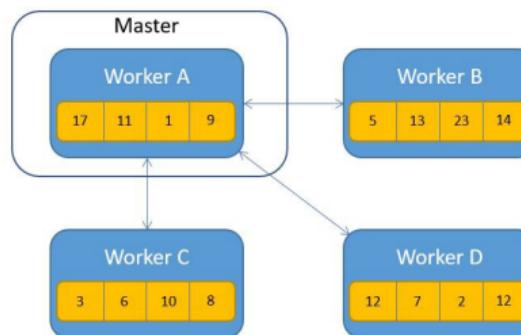
- ▶ Send the array of data to each other.
- ▶ Apply the reduction operation on each process.
- ▶ Too many unnecessary messages.



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Master-Worker AllReduce

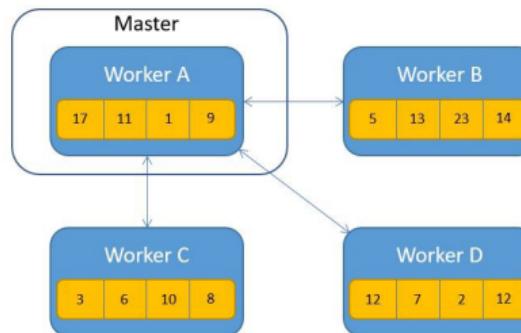
- ▶ Selecting one process as a **master**, gather all arrays into the master.
- ▶ Perform **reduction operations** locally in the **master**.
- ▶ Distribute the result to the **other processes**.



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Master-Worker AllReduce

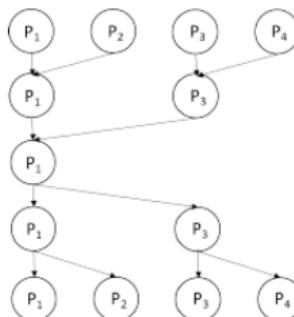
- ▶ Selecting one process as a **master**, gather all arrays into the master.
- ▶ Perform **reduction operations** locally in the **master**.
- ▶ **Distribute the result** to the **other processes**.
- ▶ The master becomes a **bottleneck** (**not scalable**).



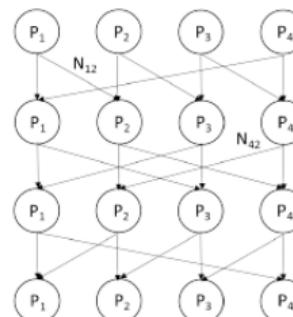
[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Other implementations

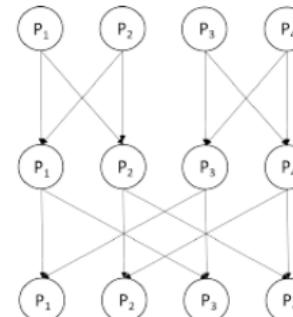
- ▶ Some try to minimize bandwidth.
- ▶ Some try to minimize latency.



(a) Tree AllReduce



(b) Round-robin AllReduce



(c) Butterfly AllReduce

[Zhao H. et al., arXiv:1312.3020, 2013]

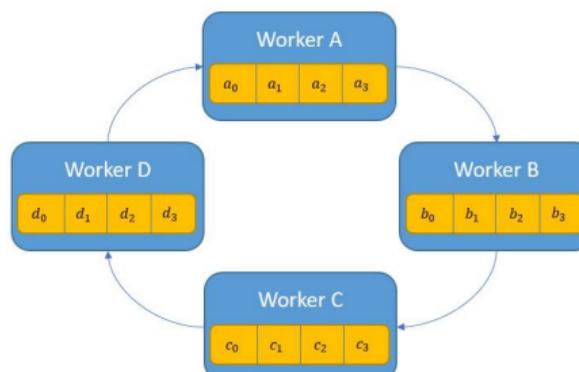


AllReduce Implementation - Ring-AllReduce (1/6)

- ▶ The **Ring-Allreduce** has two phases:
 1. First, the **share-reduce** phase
 2. Then, the **share-only** phase

AllReduce Implementation - Ring-AllReduce (2/6)

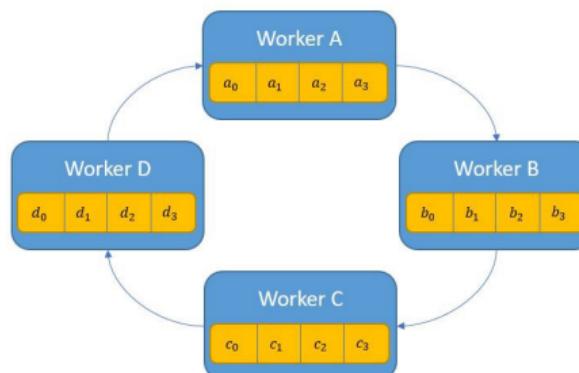
- In the **share-reduce** phase, each process p sends data to the process $(p+1) \% m$
 - m is the number of processes, and $\%$ is the modulo operator.



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Ring-AllReduce (2/6)

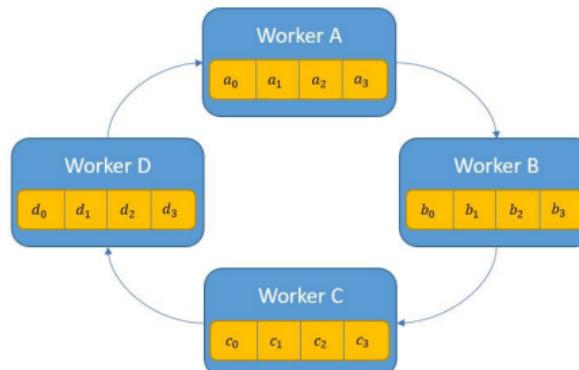
- In the **share-reduce** phase, each process p sends data to the process $(p+1) \% m$
 - m is the number of processes, and $\%$ is the modulo operator.
- The **array of data** on each process is divided to m chunks ($m=4$ here).



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Ring-AllReduce (2/6)

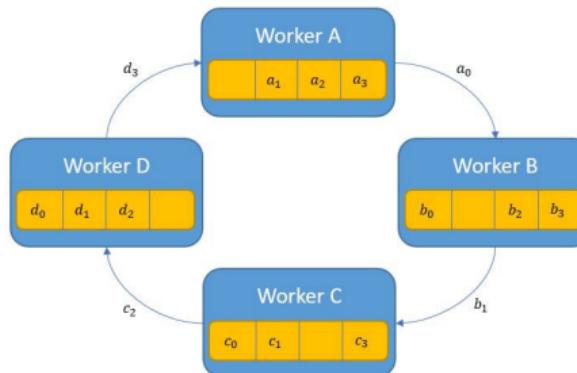
- In the **share-reduce** phase, each process p sends data to the process $(p+1) \% m$
 - m is the number of processes, and $\%$ is the modulo operator.
- The **array of data** on each process is divided to m chunks ($m=4$ here).
- Each one of these **chunks** will be **indexed** by i going forward.



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Ring-AllReduce (3/6)

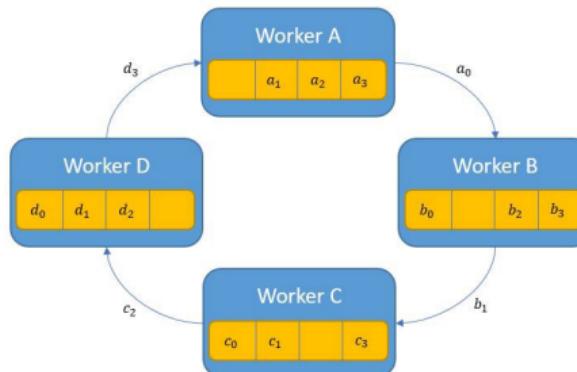
- In the first share-reduce step, process A sends a_0 to process B.



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Ring-AllReduce (3/6)

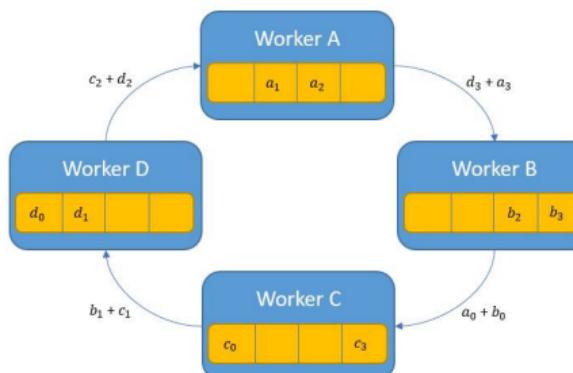
- ▶ In the **first share-reduce step**, process **A** sends **a_0** to process **B**.
- ▶ Process **B** sends **b_1** to process **C**, etc.



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Ring-AllReduce (4/6)

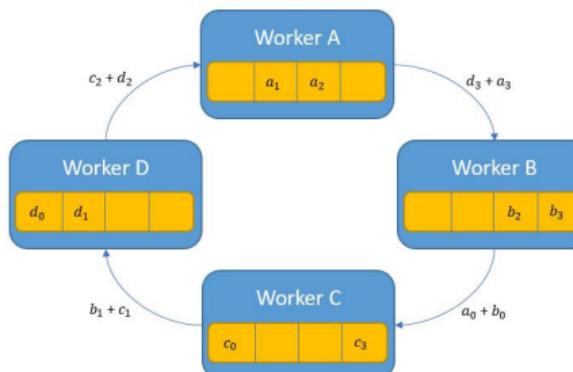
- When each process receives the data from the previous process, it applies the reduce operator (e.g., sum)



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Ring-AllReduce (4/6)

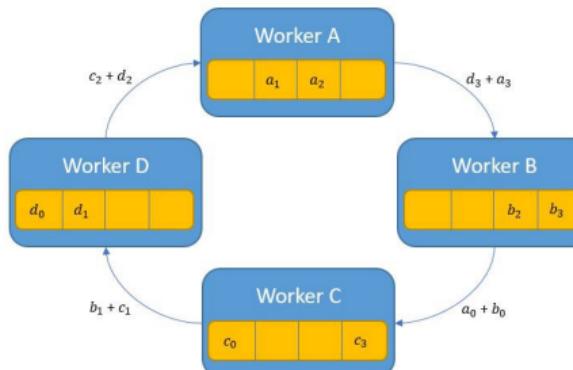
- When each process receives the data from the previous process, it applies the reduce operator (e.g., sum)
 - The reduce operator should be associative and commutative.



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Ring-AllReduce (4/6)

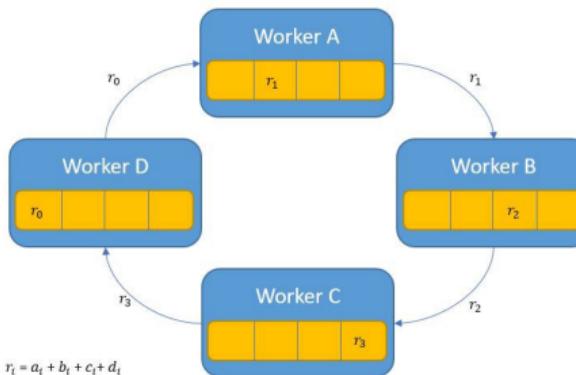
- ▶ When each process **receives** the data from the **previous process**, it applies the **reduce operator** (e.g., sum)
 - The **reduce operator** should be **associative** and **commutative**.
- ▶ It then proceeds to **send it to the next process** in the **ring**.



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Ring-AllReduce (5/6)

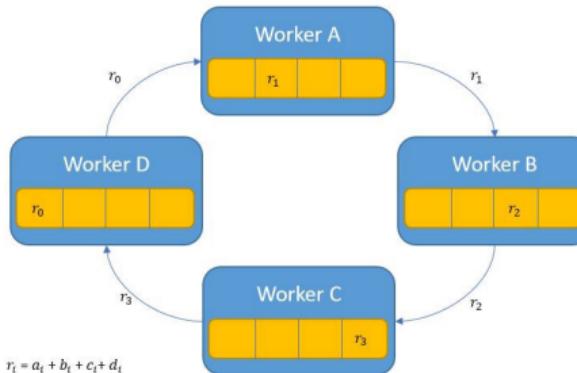
- The share-reduce phase finishes when each process holds the complete reduction of chunk i.



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Ring-AllReduce (5/6)

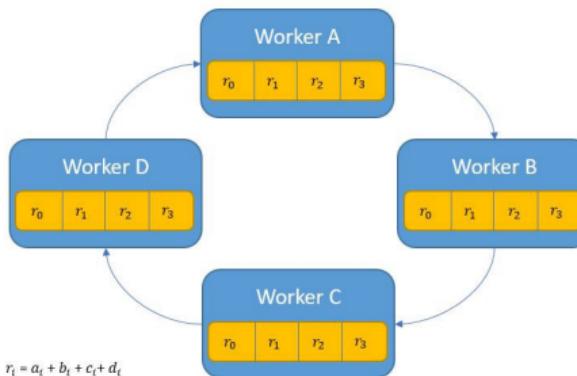
- ▶ The share-reduce phase finishes when each process holds the complete reduction of chunk i.
- ▶ At this point each process holds a part of the end result.



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Ring-AllReduce (6/6)

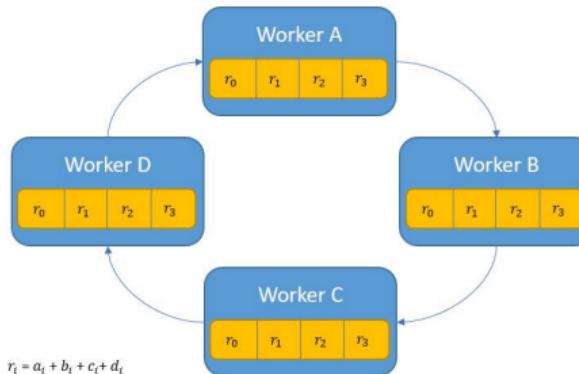
- The **share-only** step is the same process of sharing the data in a ring-like fashion without applying the reduce operation.



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]

AllReduce Implementation - Ring-AllReduce (6/6)

- ▶ The **share-only** step is the same process of sharing the data in a ring-like fashion **without applying the reduce operation**.
- ▶ This **consolidates** the **result of each chunk** in **every process**.



[<https://towardsdatascience.com/visual-intuition-on-ring-allreduce-for-distributed-deep-learning-d1f34b4911da>]



Master-Worker AllReduce vs. Ring-AllReduce

- ▶ N : number of elements, m : number of processes



Master-Worker AllReduce vs. Ring-AllReduce

- ▶ N : number of elements, m : number of processes
- ▶ Master-Worker AllReduce



Master-Worker AllReduce vs. Ring-AllReduce

- ▶ N : number of elements, m : number of processes
- ▶ Master-Worker AllReduce
 - First each process sends N elements to the master: $N \times (m - 1)$ messages.



Master-Worker AllReduce vs. Ring-AllReduce

- ▶ N : number of elements, m : number of processes
- ▶ Master-Worker AllReduce
 - First each **process** sends N elements to the **master**: $N \times (m - 1)$ messages.
 - Then the **master** sends the results back to the **process**: another $N \times (m - 1)$ messages.



Master-Worker AllReduce vs. Ring-AllReduce

- ▶ N : number of elements, m : number of processes
- ▶ Master-Worker AllReduce
 - First each process sends N elements to the master: $N \times (m - 1)$ messages.
 - Then the master sends the results back to the process: another $N \times (m - 1)$ messages.
 - Total network traffic is $2(N \times (m - 1))$, which is proportional to m .



Master-Worker AllReduce vs. Ring-AllReduce

- ▶ N : number of elements, m : number of processes
- ▶ Master-Worker AllReduce
 - First each **process** sends N elements to the **master**: $N \times (m - 1)$ messages.
 - Then the **master** sends the results back to the **process**: another $N \times (m - 1)$ messages.
 - Total network traffic is $2(N \times (m - 1))$, which is **proportional** to m .
- ▶ Ring-AllReduce



Master-Worker AllReduce vs. Ring-AllReduce

- ▶ N : number of elements, m : number of processes
- ▶ Master-Worker AllReduce
 - First each **process** sends N elements to the **master**: $N \times (m - 1)$ messages.
 - Then the **master** sends the results back to the **process**: another $N \times (m - 1)$ messages.
 - Total network traffic is $2(N \times (m - 1))$, which is **proportional** to m .
- ▶ Ring-AllReduce
 - In the **share-reduce** step each **process** sends $\frac{N}{m}$ elements, and it does it $m - 1$ times:
 $\frac{N}{m} \times (m - 1)$ messages.

Master-Worker AllReduce vs. Ring-AllReduce

- ▶ N : number of elements, m : number of processes
- ▶ Master-Worker AllReduce
 - First each **process** sends N elements to the **master**: $N \times (m - 1)$ messages.
 - Then the **master** sends the results back to the **process**: another $N \times (m - 1)$ messages.
 - Total network traffic is $2(N \times (m - 1))$, which is **proportional** to m .
- ▶ Ring-AllReduce
 - In the **share-reduce** step each **process** sends $\frac{N}{m}$ elements, and it does it $m - 1$ times:
 $\frac{N}{m} \times (m - 1)$ messages.
 - On the **share-only** step, each **process** sends the result for the chunk it calculated: another
 $\frac{N}{m} \times (m - 1)$ messages.

Master-Worker AllReduce vs. Ring-AllReduce

- ▶ N : number of elements, m : number of processes
- ▶ Master-Worker AllReduce
 - First each **process** sends N elements to the **master**: $N \times (m - 1)$ messages.
 - Then the **master** sends the results back to the **process**: another $N \times (m - 1)$ messages.
 - Total network traffic is $2(N \times (m - 1))$, which is **proportional** to m .
- ▶ Ring-AllReduce
 - In the **share-reduce** step each **process** sends $\frac{N}{m}$ elements, and it does it $m - 1$ times:
 $\frac{N}{m} \times (m - 1)$ messages.
 - On the **share-only** step, each **process** sends the result for the chunk it calculated: another
 $\frac{N}{m} \times (m - 1)$ messages.
 - Total network traffic is $2(\frac{N}{m} \times (m - 1))$.



Communication Synchronization and Frequency



Synchronization

- ▶ When to synchronize the parameters among the parallel workers?



Communication Synchronization (1/2)

- ▶ Synchronizing the model replicas in **data-parallel** training requires **communication**
 - between **workers**, in **allreduce**
 - between **workers and parameter servers**, in the **centralized architecture**



Communication Synchronization (1/2)

- ▶ Synchronizing the model replicas in data-parallel training requires communication
 - between workers, in allreduce
 - between workers and parameter servers, in the centralized architecture
- ▶ The communication synchronization decides how frequently all local models are synchronized with others.



Communication Synchronization (2/2)

- ▶ It will influence:
 - The communication **traffic**
 - The **performance**
 - The **convergence** of model training



Communication Synchronization (2/2)

- ▶ It will influence:
 - The communication **traffic**
 - The **performance**
 - The **convergence** of model training
- ▶ There is a **trade-off** between the communication **traffic** and the **convergence**.



Reducing Synchronization Overhead

- ▶ Two directions for improvement:



Reducing Synchronization Overhead

- ▶ Two directions for improvement:
 1. To **relax** the **synchronization** among all workers.



Reducing Synchronization Overhead

- ▶ Two directions for improvement:
 1. To **relax** the **synchronization** among all workers.
 2. The **frequency of communication** can be **reduced** by more computation in one iteration.

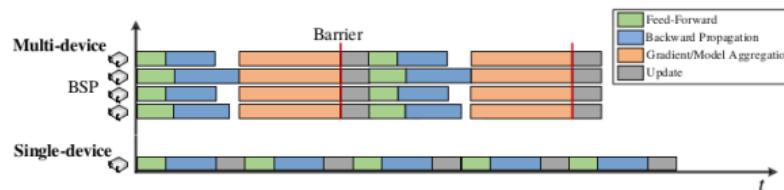


Communication Synchronization Models

- ▶ Synchronous
- ▶ Stale-synchronous
- ▶ Asynchronous
- ▶ Local SGD

Communication Synchronization - Synchronous

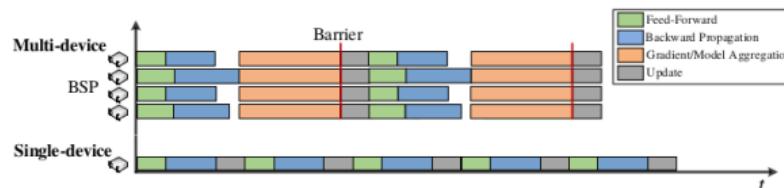
- ▶ After each **iteration**, the workers **synchronize** their parameter updates.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Synchronous

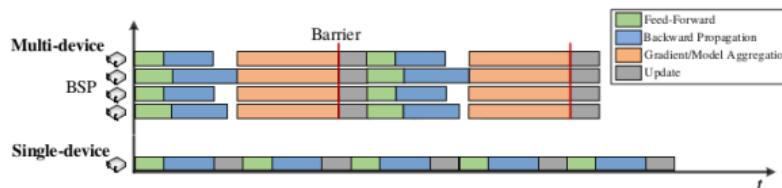
- ▶ After each **iteration**, the workers **synchronize** their parameter updates.
- ▶ Every worker must **wait** for **all workers** to **finish** the transmission of all parameters in the current iteration, before the **next training**.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Synchronous

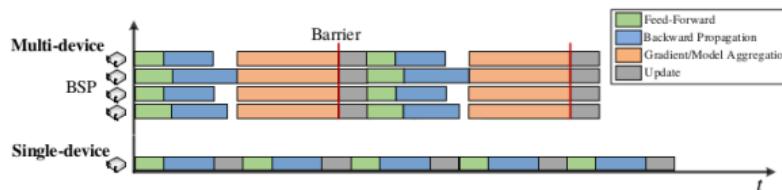
- ▶ After each **iteration**, the workers **synchronize** their parameter updates.
- ▶ Every worker must **wait** for **all workers** to **finish** the transmission of all parameters in the current iteration, before the **next training**.
- ▶ **Stragglers** can influence the overall system **throughput**.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Synchronous

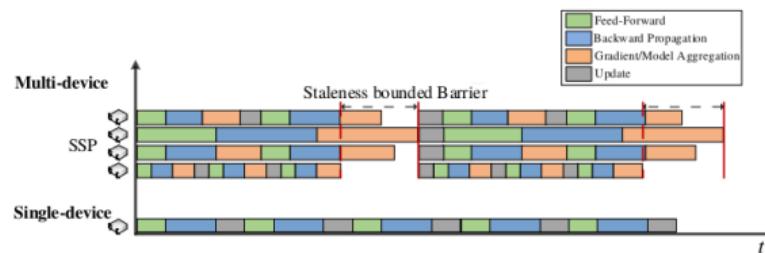
- ▶ After each **iteration**, the workers **synchronize** their parameter updates.
- ▶ Every worker must **wait** for **all workers** to **finish** the transmission of all parameters in the current iteration, before the **next training**.
- ▶ **Stragglers** can influence the overall system **throughput**.
- ▶ High **communication** cost that **limits** the system **scalability**.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Stale Synchronous (1/2)

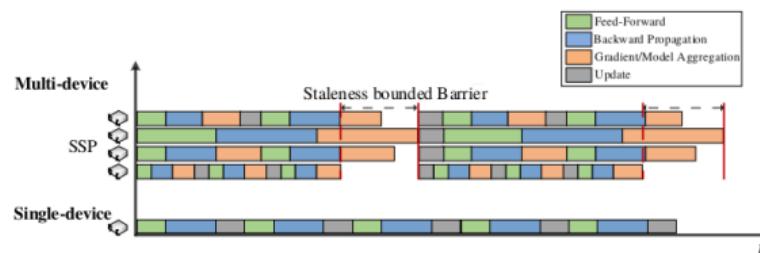
- ▶ Alleviate the straggler problem without losing synchronization.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Stale Synchronous (1/2)

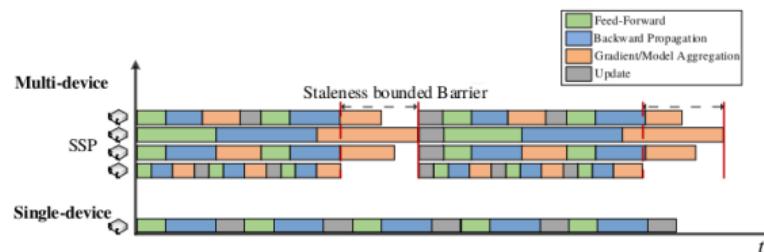
- ▶ Alleviate the straggler problem without losing synchronization.
- ▶ The faster workers to do **more updates** than the **slower workers** to **reduce the waiting time** of the faster workers.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Stale Synchronous (1/2)

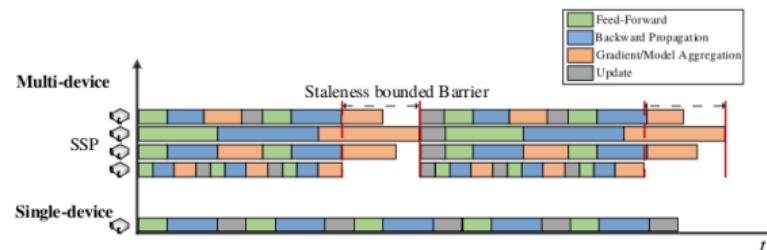
- ▶ Alleviate the straggler problem without losing synchronization.
- ▶ The faster workers to do **more updates** than the **slower workers** to reduce the waiting time of the faster workers.
- ▶ **Staleness bounded barrier** to limit the **iteration gap** between the fastest worker and the slowest worker.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Stale Synchronous (2/2)

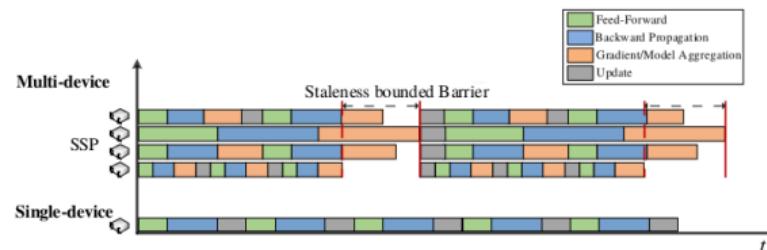
- ▶ For a maximum staleness bound s , the update formula of worker i at iteration $t+1$:
- ▶ $w_{i,t+1} := w_0 - \eta(\sum_{k=1}^t \sum_{j=1}^n G_{j,k} + \sum_{k=t-s}^t G_{i,k} + \sum_{(j,k) \in S_{i,t+1}} G_{j,k})$



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Stale Synchronous (2/2)

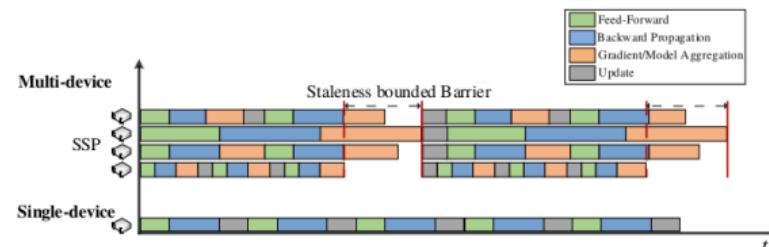
- ▶ For a maximum staleness bound s , the update formula of worker i at iteration $t+1$:
- ▶ $w_{i,t+1} := w_0 - \eta(\sum_{k=1}^t \sum_{j=1}^n G_{j,k} + \sum_{k=t-s}^t G_{i,k} + \sum_{(j,k) \in S_{i,t+1}} G_{j,k})$
- ▶ The update has three parts:



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Stale Synchronous (2/2)

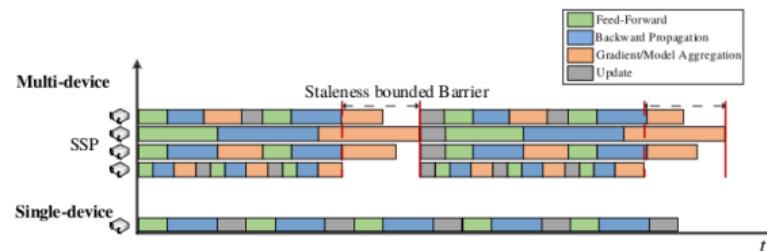
- ▶ For a maximum staleness bound s , the update formula of worker i at iteration $t+1$:
- ▶ $w_{i,t+1} := w_0 - \eta(\sum_{k=1}^t \sum_{j=1}^n G_{j,k} + \sum_{k=t-s}^t G_{i,k} + \sum_{(j,k) \in S_{i,t+1}} G_{j,k})$
- ▶ The update has three parts:
 1. Guaranteed pre-window updates from clock 1 to t over all workers.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Stale Synchronous (2/2)

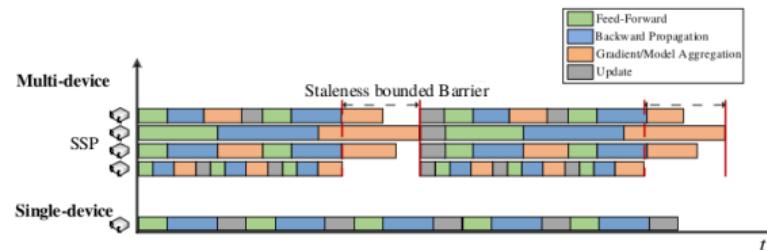
- ▶ For a maximum staleness bound s , the update formula of worker i at iteration $t+1$:
- ▶ $w_{i,t+1} := w_0 - \eta(\sum_{k=1}^t \sum_{j=1}^n G_{j,k} + \sum_{k=t-s}^t G_{i,k} + \sum_{(j,k) \in S_{i,t+1}} G_{j,k})$
- ▶ The update has three parts:
 1. Guaranteed pre-window updates from clock 1 to t over all workers.
 2. Guaranteed read-my-writes in-window updates made by the querying worker i .



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Stale Synchronous (2/2)

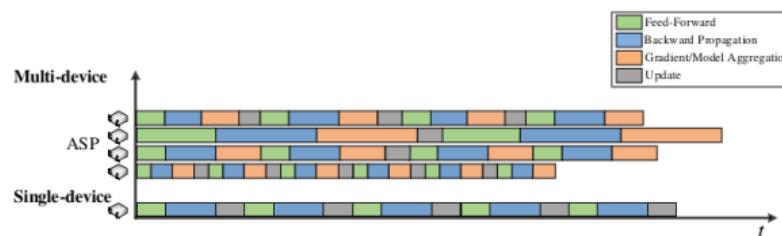
- ▶ For a maximum staleness bound s , the update formula of worker i at iteration $t+1$:
- ▶ $w_{i,t+1} := w_0 - \eta(\sum_{k=1}^t \sum_{j=1}^n G_{j,k} + \sum_{k=t-s}^t G_{i,k} + \sum_{(j,k) \in S_{i,t+1}} G_{j,k})$
- ▶ The update has three parts:
 1. Guaranteed pre-window updates from clock 1 to t over all workers.
 2. Guaranteed read-my-writes in-window updates made by the querying worker i .
 3. Best-effort in-window updates. $S_{i,t+1}$ is some subset of the updates from other workers during period $[t - s]$.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Asynchronous (1/2)

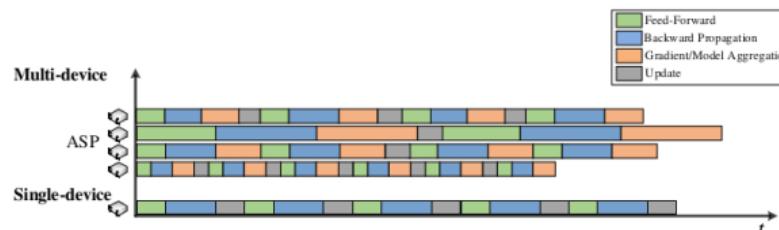
- ▶ It completely eliminates the synchronization.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Asynchronous (1/2)

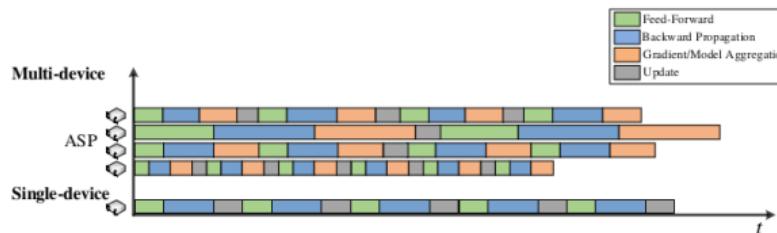
- ▶ It completely **eliminates** the synchronization.
- ▶ Each work **transmits its gradients** to the PS after it calculates the gradients.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Asynchronous (1/2)

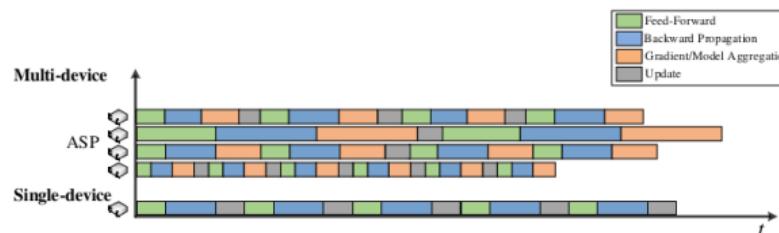
- ▶ It completely **eliminates** the synchronization.
- ▶ Each work **transmits its gradients** to the PS after it calculates the gradients.
- ▶ The PS updates the global model **without waiting** for the other workers.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Asynchronous (2/2)

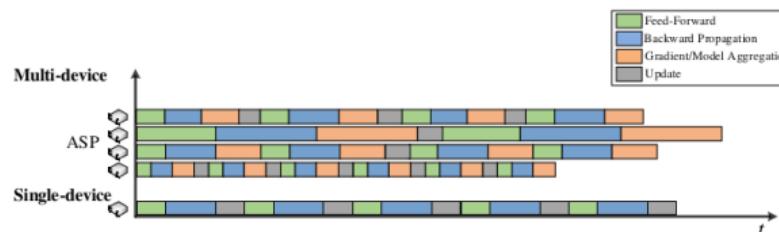
► $w_{t+1} := w_t - \eta \sum_{i=1}^n G_{i,t-\tau_{k,i}}$



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Asynchronous (2/2)

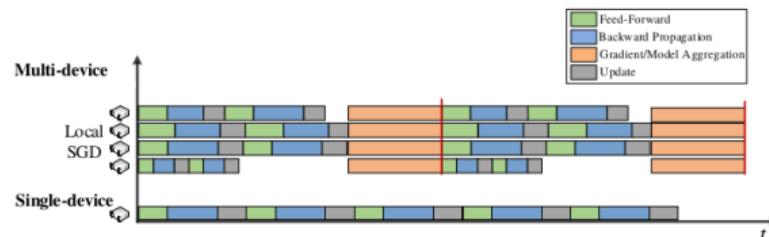
- ▶ $w_{t+1} := w_t - \eta \sum_{i=1}^n G_{i,t-\tau_{k,i}}$
- ▶ $\tau_{k,i}$ is the time delay between the moment when worker i calculates the gradient at the current iteration.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Local SGD

- ▶ All workers **run several iterations**, and then **averages all local models** into the newest global model.

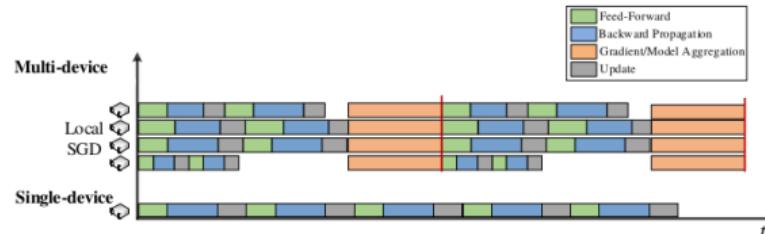


[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Synchronization - Local SGD

- ▶ All workers **run several iterations**, and then **averages all local models** into the newest global model.
- ▶ If \mathcal{I}_T represents the synchronization timestamps, then:

$$w_{i,t+1} = \begin{cases} w_{i,t} - \eta g_{i,t} & \text{if } t+1 \notin \mathcal{I}_T \\ w_{i,t} - \eta \frac{1}{n} \sum_{i=1}^n g_{i,t} & \text{if } t+1 \in \mathcal{I}_T \end{cases}$$



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]



Communication Compression



Communication Compression

- ▶ Reduce the communication traffic with **little impact** on the model convergence.



Communication Compression

- ▶ Reduce the communication traffic with **little impact** on the model convergence.
- ▶ Compress the exchanged gradients or models **before transmitting** across the network.



Communication Compression

- ▶ Reduce the communication traffic with **little impact** on the model convergence.
- ▶ Compress the exchanged gradients or models **before transmitting** across the network.
- ▶ Quantization

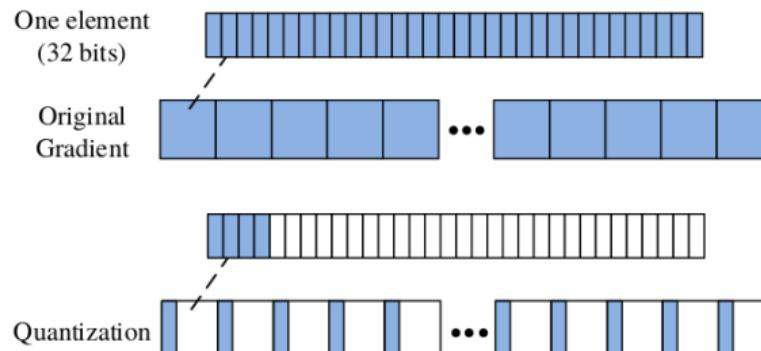


Communication Compression

- ▶ Reduce the communication traffic with **little impact** on the model convergence.
- ▶ Compress the exchanged gradients or models **before transmitting** across the network.
- ▶ Quantization
- ▶ Sparsification

Communication Compression - Quantization

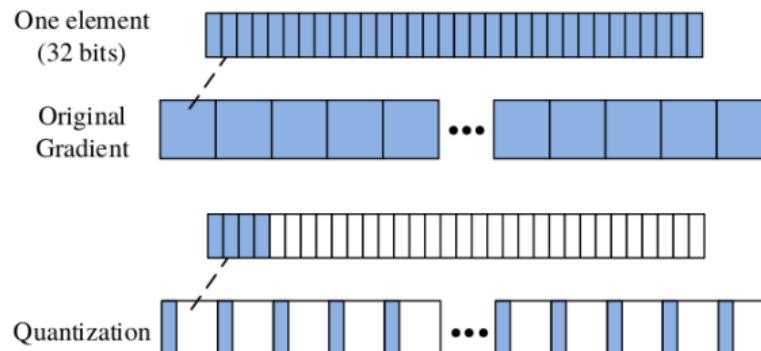
- ▶ Using lower bits to represent the data.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Compression - Quantization

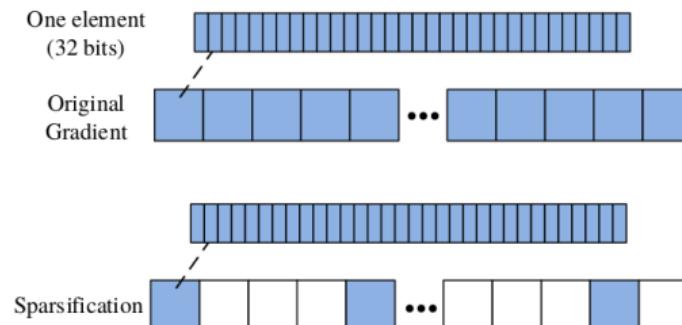
- ▶ Using lower bits to represent the data.
- ▶ The gradients are of low precision.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Compression - Sparsification

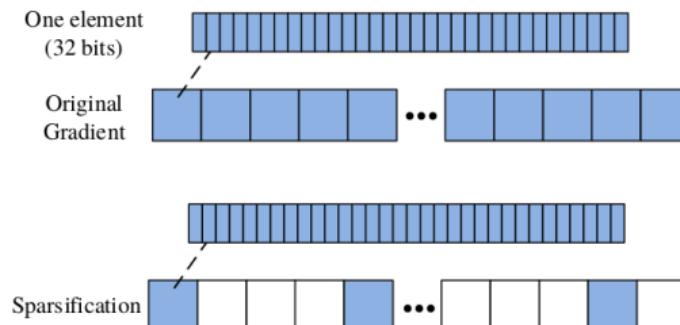
- ▶ Reducing the **number of elements** that are transmitted at each iteration.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Compression - Sparsification

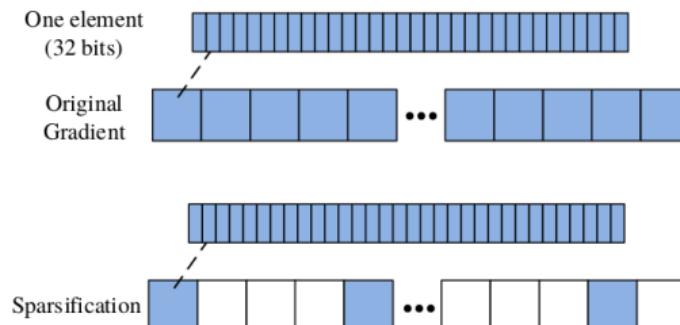
- ▶ Reducing the **number of elements** that are transmitted at each iteration.
- ▶ Only **significant gradients** are required to **update the model parameter** to guarantee the convergence of the training.



[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

Communication Compression - Sparsification

- ▶ Reducing the **number of elements** that are transmitted at each iteration.
- ▶ Only **significant gradients** are required to **update the model parameter** to guarantee the convergence of the training.
- ▶ E.g., the **zero-valued** elements are no need to transmit.



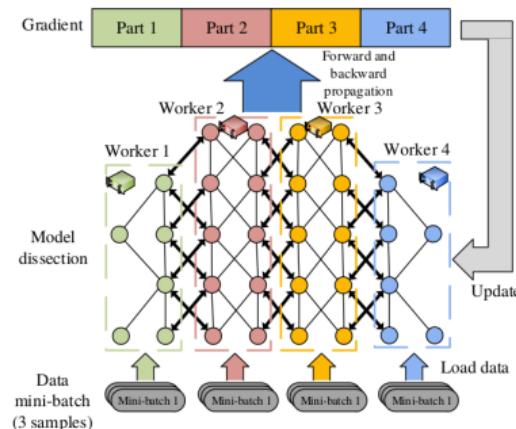
[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]



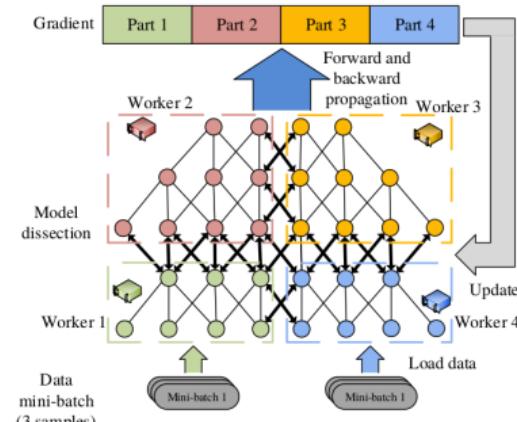
Model Parallelism

Model Parallelization

- The model is split across multiple devices.

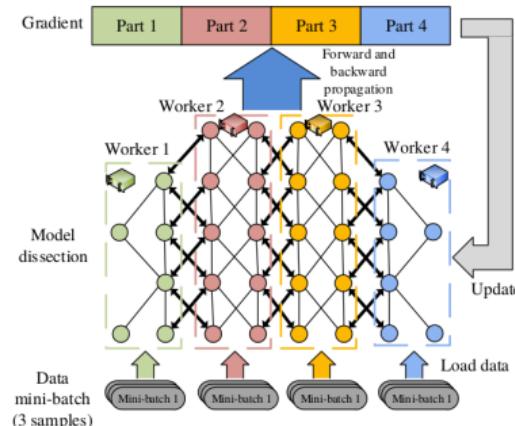


[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]

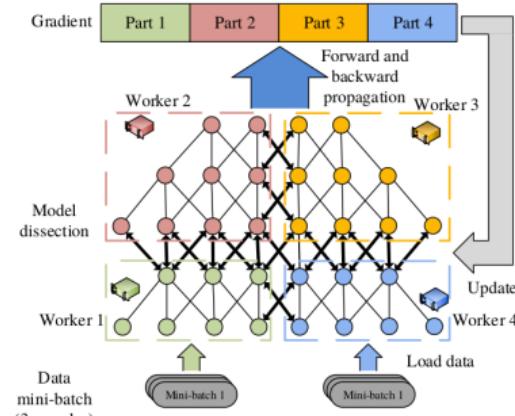


Model Parallelization

- The model is split across multiple devices.
- Depends on the architecture of the NN.

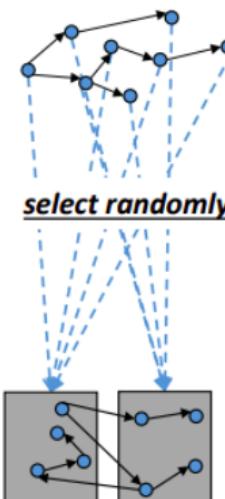


[Tang et al., Communication-Efficient Distributed Deep Learning: A Comprehensive Survey, 2020]



Model Parallelization - Hash Partitioning

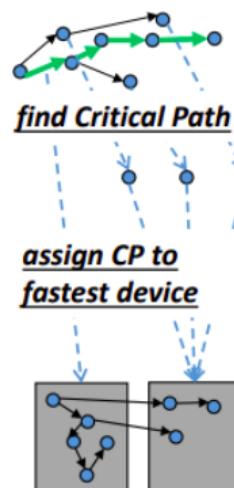
- ▶ Randomly assign vertices to devices proportionally to the capacity of the devices by using a **hash function**.



[Mayer, R. et al., The TensorFlow Partitioning and Scheduling Problem, 2017]

Model Parallelization - Critical Path

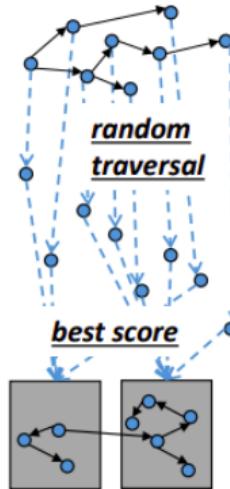
- ▶ Assigning the complete **critical path** to the fastest device.
- ▶ **Critical path**: the path with the **longest computation time** from source to sink vertex.



[Mayer, R. et al., The TensorFlow Partitioning and Scheduling Problem, 2017]

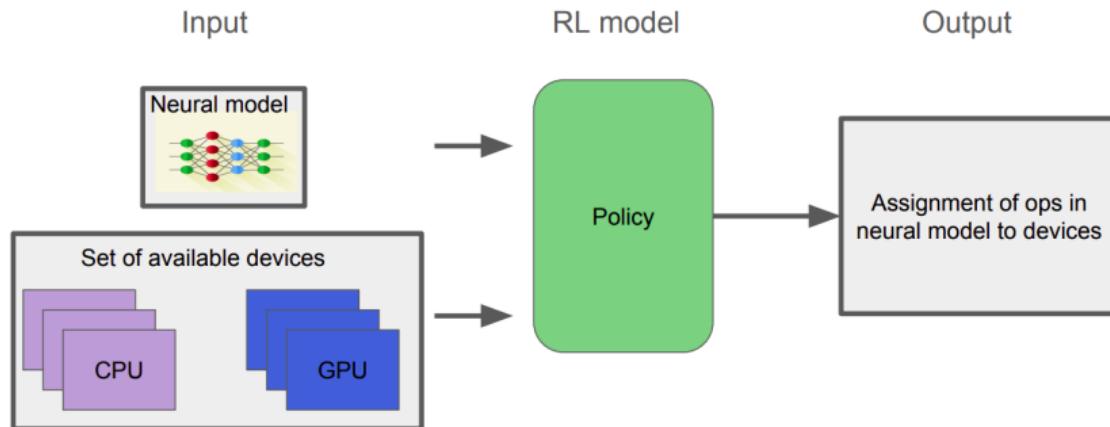
Model Parallelization - Multi-Objective Heuristics

- ▶ Different **objectives**, e.g., memory, importance, traffic, and execution time



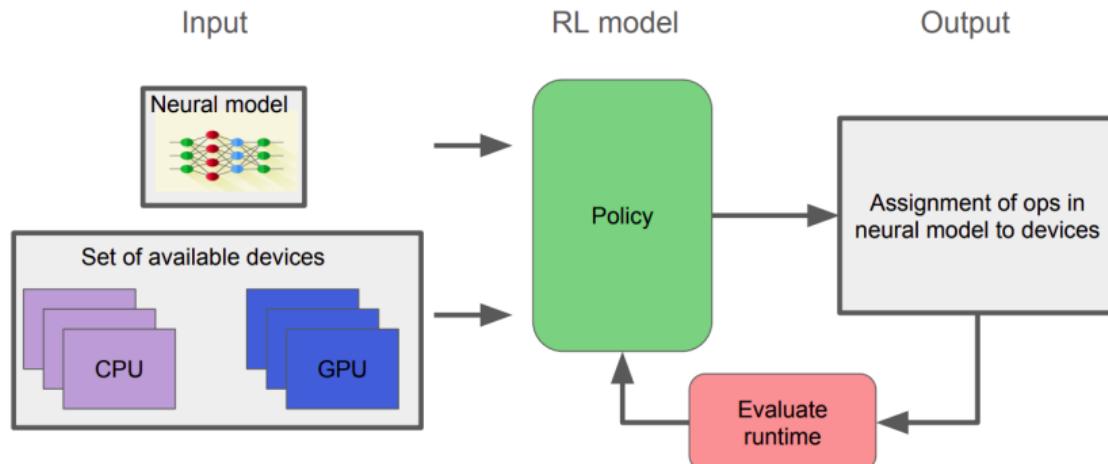
[Mayer, R. et al., The TensorFlow Partitioning and Scheduling Problem, 2017]

Model Parallelization - Reinforcement Learning (1/5)



[Mirhoseini et al., Device Placement Optimization with Reinforcement Learning, 2017]

Model Parallelization - Reinforcement Learning (2/5)



[Mirhoseini et al., Device Placement Optimization with Reinforcement Learning, 2017]



Model Parallelization - Reinforcement Learning (3/5)

- ▶ $J(w) = \mathbb{E}_{\mathcal{P} \sim \pi(\mathcal{P}|\mathcal{G}, w)}[R(\mathcal{P})|\mathcal{G}]$
- ▶ Objective: $\arg \min_w J(w)$



Model Parallelization - Reinforcement Learning (3/5)

- ▶ $J(w) = \mathbb{E}_{\mathcal{P} \sim \pi(\mathcal{P}|\mathcal{G}, w)}[R(\mathcal{P})|\mathcal{G}]$
- ▶ Objective: $\arg \min_w J(w)$
- ▶ \mathcal{G} : input neural graph



Model Parallelization - Reinforcement Learning (3/5)

- ▶ $J(w) = \mathbb{E}_{\mathcal{P} \sim \pi(\mathcal{P}|\mathcal{G}, w)}[R(\mathcal{P})|\mathcal{G}]$
- ▶ Objective: $\arg \min_w J(w)$
- ▶ \mathcal{G} : input neural graph
- ▶ R : runtime

Model Parallelization - Reinforcement Learning (3/5)

- ▶ $J(w) = \mathbb{E}_{\mathcal{P} \sim \pi(\mathcal{P}|\mathcal{G}, w)}[R(\mathcal{P})|\mathcal{G}]$
- ▶ Objective: $\arg \min_w J(w)$
- ▶ \mathcal{G} : input neural graph
- ▶ R : runtime
- ▶ $J(w)$: expected runtime

Model Parallelization - Reinforcement Learning (3/5)

- ▶ $J(w) = \mathbb{E}_{\mathcal{P} \sim \pi(\mathcal{P}|\mathcal{G}, w)}[R(\mathcal{P})|\mathcal{G}]$
- ▶ Objective: $\arg \min_w J(w)$
- ▶ \mathcal{G} : input neural graph
- ▶ R : runtime
- ▶ $J(w)$: expected runtime
- ▶ w : trainable parameters of policy



Model Parallelization - Reinforcement Learning (3/5)

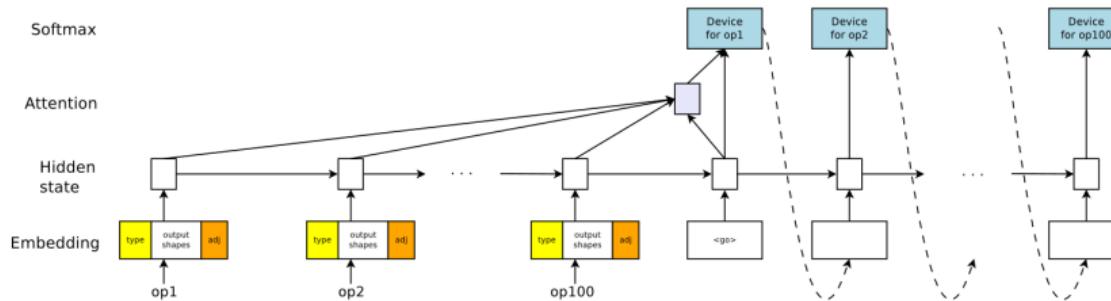
- ▶ $J(w) = \mathbb{E}_{\mathcal{P} \sim \pi(\mathcal{P}|\mathcal{G}, w)}[R(\mathcal{P})|\mathcal{G}]$
- ▶ Objective: $\arg \min_w J(w)$
- ▶ \mathcal{G} : input neural graph
- ▶ R : runtime
- ▶ $J(w)$: expected runtime
- ▶ w : trainable parameters of policy
- ▶ $\pi(\mathcal{P}|\mathcal{G}, w)$: policy

Model Parallelization - Reinforcement Learning (3/5)

- ▶ $J(w) = \mathbb{E}_{\mathcal{P} \sim \pi(\mathcal{P}|\mathcal{G}, w)}[R(\mathcal{P})|\mathcal{G}]$
- ▶ Objective: $\arg \min_w J(w)$
- ▶ \mathcal{G} : input neural graph
- ▶ R : runtime
- ▶ $J(w)$: expected runtime
- ▶ w : trainable parameters of policy
- ▶ $\pi(\mathcal{P}|\mathcal{G}, w)$: policy
- ▶ \mathcal{P} : output placements $\in \{1, 2, \dots, num_ops\}^{num_devices}$

Model Parallelization - Reinforcement Learning (4/5)

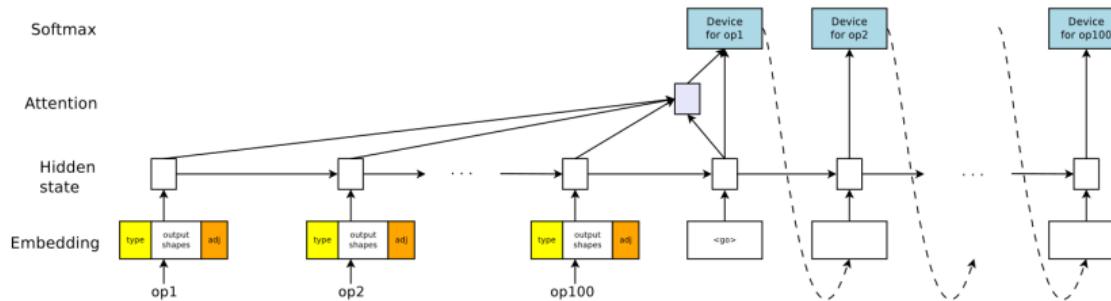
- ▶ RL reward function based on execution runtime.



[Mirhoseini et al., Device Placement Optimization with Reinforcement Learning, 2017]

Model Parallelization - Reinforcement Learning (4/5)

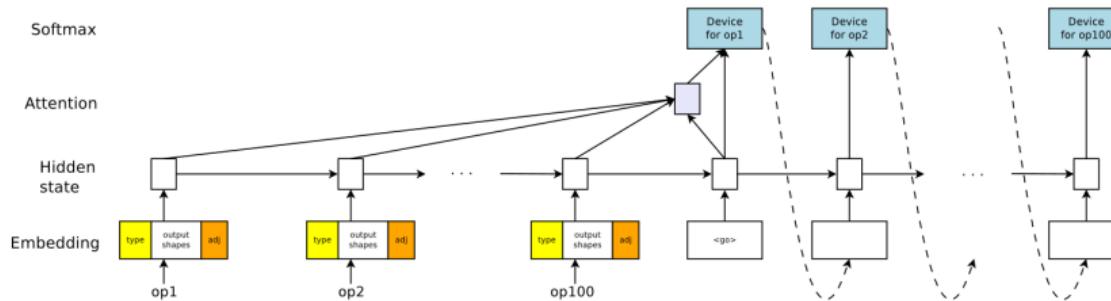
- ▶ RL reward function based on execution runtime.
- ▶ The RL policy is defined as a seq-to-seq model.



[Mirhoseini et al., Device Placement Optimization with Reinforcement Learning, 2017]

Model Parallelization - Reinforcement Learning (4/5)

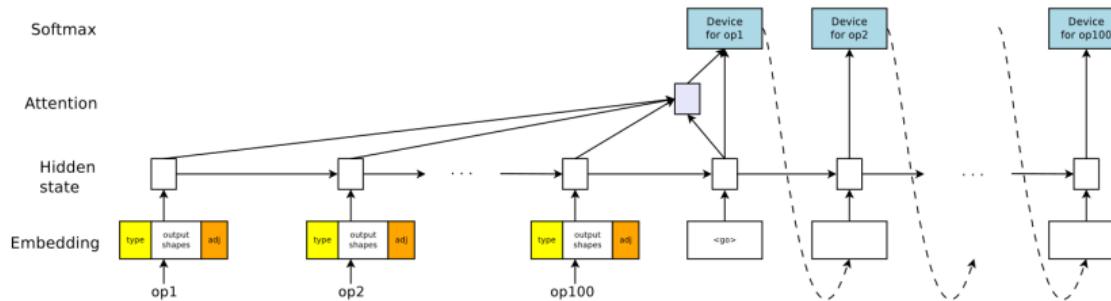
- ▶ RL reward function based on execution runtime.
- ▶ The RL policy is defined as a seq-to-seq model.
- ▶ RNN Encoder receives graph embedding for each operation.



[Mirhoseini et al., Device Placement Optimization with Reinforcement Learning, 2017]

Model Parallelization - Reinforcement Learning (4/5)

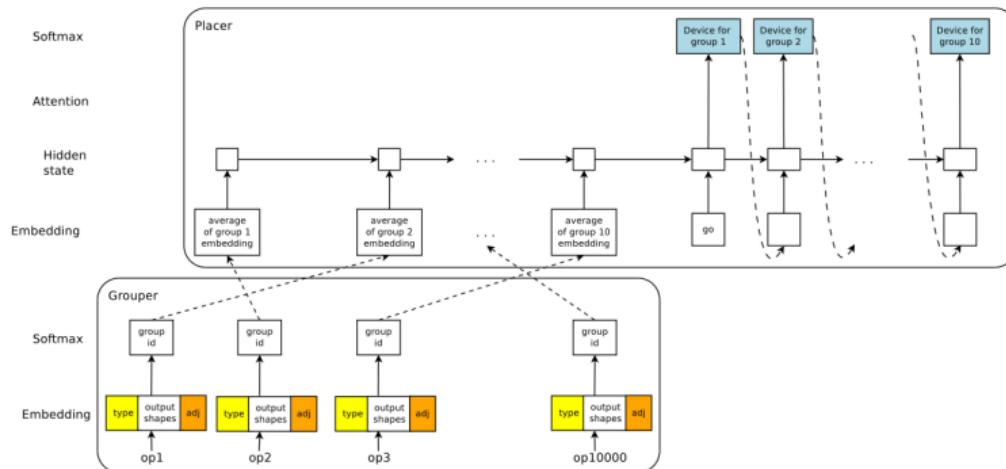
- ▶ RL reward function based on execution runtime.
- ▶ The RL policy is defined as a seq-to-seq model.
- ▶ RNN Encoder receives graph embedding for each operation.
- ▶ RNN Decoder predicts a device placement for each operation.



[Mirhoseini et al., Device Placement Optimization with Reinforcement Learning, 2017]

Model Parallelization - Reinforcement Learning (5/5)

- ▶ Grouping operations.
- ▶ Prediction is for **group placement**, not for a single operation.



[Mirhoseini et al., A Hierarchical Model for Device Placement, 2018]



Summary



Summary

- ▶ Scalability matters
- ▶ Parallelization
- ▶ Data Parallelization
 - Parameter server vs. AllReduce
 - Synchronized vs. asynchronous
- ▶ Model Parallelization
 - Random, critical path, multi-objective, RL



Thanks!